


# NGS data: Beauty and the Beast

*How to infer population genetics parameters from messy sequencing data*

Matteo Fumagalli

12<sup>th</sup> September 2017

## Goal of the day





## Presentation outline

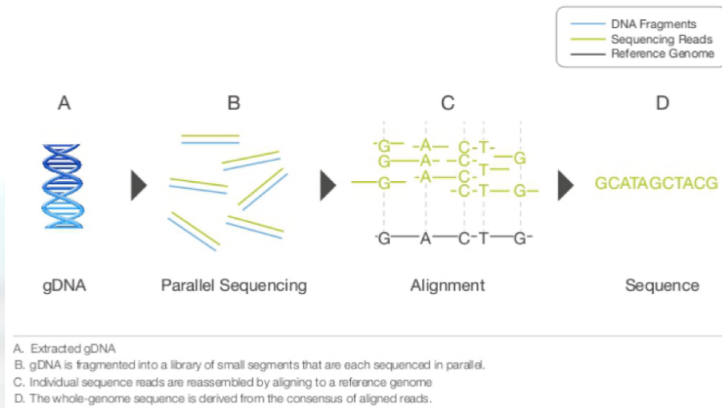
- 1 Introduction
- 2 Genotype likelihoods
- 3 Genotype calling, really?  
Low-depth (genotypes)
- 4 Allele frequencies  
Low-depth (allele frequencies)
- 5 Experimental design

# Sanger sequencing

aka first/former generation sequencing



# Next Generation Sequencing



## Low-level data

## FASTQ

```
a'X_\Va\J'KaYJHG^]b\a^BBBBBBBBBBBBBB <-- quality score
@FC42BF1AAXX:6:1:5:732#0/1 <-- read ID
TGATTCTCTCGATATCCAGTCCTTAGTGNCATAGN <-- read (bases)
+
a^_aaaa'aa'_aaa_aaa'__'_'VBBBBBBBBB
@FC42BF1AAXX:6:1:5:492#0/1
AACAGTGGGAGGCTGCAGCAGGAGGATTNCTGAAN
+
ababb_abbbaab^'aaTaabbaBBBBBBBBB
@FC42BF1AAXX:6:1:5:480#0/1
ACCTCCTCAGAGTTCTCGAGCTCGAGAANTCTGGN
```

## Quality scores

### Qscore

- The ASCII values can be interpreted as a probability
- A Q20 (ASCII 'T') score is probability of 1%
- The score is the probability,  $P$ , that the base is incorrect
- 

$$Q_{score} = -10\log_{10}(P)$$

- 

$$P = 10^{\frac{-Q}{10}}$$

## Quality scores

The qscores are encoded as ASCII characters, and are shifted by +33 (now the standard) or +64.

Phred Quality Score	Probability of error	Base call accuracy
0 ... 9	1 ... 0.13	!"#\$%&'()*
10 ... 19	0.1 ... 0.013	+,-./01234
20 ... 29	0.01 ... 0.0013	56789;:<=>
30 ... 39	0.001 ... 0.00013	?@ABCDEFGH
40 ... 49	0.0001 ... 0.000014	IJKLMNOPQR



## Quality scores

### Example

From a fastq file we observe an 'A' with a qscore encoded as '7'. What is the probability of being 'A'? What is the probability of being 'G'?

1. We find that the corresponding ASCII value of '7' is ?  
(hint: <http://www.asciitable.com>)

## Quality scores

### Example

From a fastq file we observe an 'A' with a qscore encoded as '7'. What is the probability of being 'A'? What is the probability of being 'G'?


1. We find that the corresponding ASCII value of '7' is ?  
(hint: <http://www.asciitable.com>)
2. We subtract 33 to get a value of ?. This is our qscore.
3. The probability of 'A' being incorrect is ? (hint:  $p = 10^{\frac{-Q}{10}}$ )
4. The probability of 'A' being correct is ?
5. The probability of being 'G' (or 'C' or 'T') is ?

## FastQ files

```
@HWI-ST397_0000:2:1:2248:2126#CTTGTA/1
TTGGATCTGAAAGATGAATGTGAGAGACACAATCCAAGTCATCTCTC
ATG
+HWI-ST397_0000:2:1:2248:2126#CTTGTA/1
eeee\dZddaddddddeeeeeeeedaed_ec_ab_\NSRNRcdddc[_c^d
```

- sequencer
- flowcell
- lane/cell/tile
- position within the tile
- barcode id for pooling/multiplexing
- pair
- ...

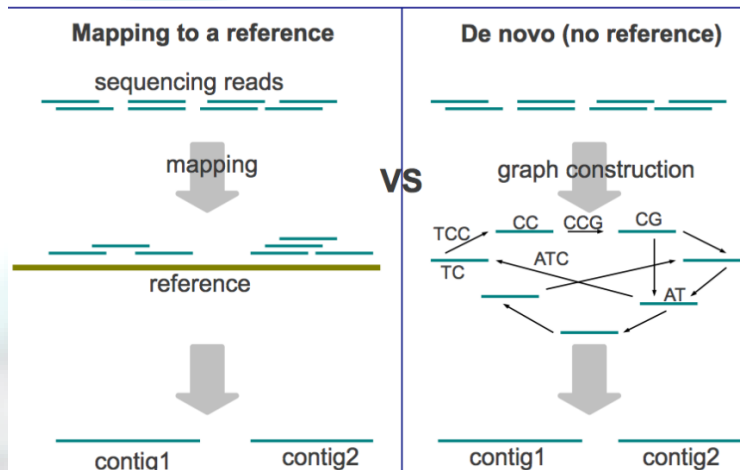
# Quality check of fastq files: fastQC



- distribution of qscores over read
- overrepresented kmers
- ...

<http://www.bioinformatics.babraham.ac.uk/projects/fastqc/>

# Alignment of reads



## Mapping to a reference

Issues:

- (i) Millions of short reads.
- (ii) Blat/blast is too slow.

But:

Mate-pair gives additional information.

Many aligners exists:

- Soap/soap2 (BGI)
- Maq/bwa (Heng Li)
- Bowtie/bowtie2 (Langmead B)
- Eland, SSAHA2, RMAP, shrimp, zoom, GEM, snap, novoAlign.

Most are based on burrows wheeler transform BTW  
(BWA, Bowtie, Soap2, ...)

## Mapping to a reference

Many different aligners, what's the difference?

- Memory usage
- Speed
- Gapped (indels)
- Using qscores (bwa pssm)
- Estimating a **mappability score**
- Multiple best hits
- Paired end data
- Output **formats** (SAM,...)

## Alignment file

an alignment file includes

reads TTTGTTCTTTCTTTCTCTCTAGTCTTCTT ...

Qscore NVFVN]^['^\_]^^U]] '['[\_VS[\_^Z]\_ ...

start position chr4 53351385

multiple best hits 1

Number of mismatch 2

sequence strand -

read quality\* V



>ARPM2ref|NC\_000001.10|:2938046-2939467 homo sapiens chromosome 1, GRCh37 primary reference assembly


TGGAAGAGGCTCAGCAGGCCAGGCCACCTGGAGGGAGAGCAGACCTGCGGCTGAGGATGCAGGGCTCC  
 CGGGCACGGTGCTAGCCCTGCCTTGAGACACCCGAGAGCTGTGGGAAGAGCTGTGGGATCCCCATTATGC  
 ATCACAAAGCGGCCCTTGAGGGCTGGTCTTTATTTTGATGAGGCTGAGAAGGGAAGGCTGCGGGCATGTT  
 TAATCCGACAGCTTTAGACTCCCGGGCTGTGATTTTGCAGATGGCTCGGGGTCTGCAAAGCGGGCTG  
 TCTGGGGAGTTTGGAGCCCCGACATGGTGAGCTCCATCTGGGGGACCTGAAGAAATCAGGCTCCCTCAG

CCAATGATTTTTTCCGTTGTTTCAGAATACGGTTAA  
+SRR038845.1 HWI-EAS038:6:1:0:1474 length=36  
BCCBA@B@B@BBBBAB@B9B@=BABA@A:@693:@B=  
@SRR038845.53 HWI-EAS038:6:1:1:360 length=36  
GTTCAAAAAGACTAAATTGTGTCAATGAAAACTC  
+SRR038845.53 HWI-EAS038:6:1:1:360 length=36

[illegible]

```
##fileformat=VCFv4.
##fileDate=20140930
##source=23andme2vcf.pl https://github.com/arrogantrobot/23andme2vcf
##reference=file:///23andme_v3_hg19_ref.txt.gz
##FORMAT=
##CHROM POS ID REF ALT QUAL FILTER INFO FORMAT GENOTYPE
chr1 81254 rs4477212 A . . . . GT
/0
chr1 752566 rs3094315 G A . . . . GT
/1
chr1 752721 rs3131972 A G . . . . GT
/1
chr1 798959 rs11240777 G . . . . GT
/0
chr1 800007 rs6681049 T C . . . . GT
/1
```

# Our data: mapped reads with quality scores



- Coverage: fraction of the genome with data
- Depth: number of reads mapped to a position
- Counts: number of different alleles mapped to a position
- Effective Base Depth: similar to the counts, but weighing for qscores and mapping quality

## Challenges




- Variable and low depth
- High sequencing and mapping errors



Quality control filters

# The data



- is a **nucleotide**/base/allele with a certain **quality** score

## Genotype likelihoods

### Likelihood

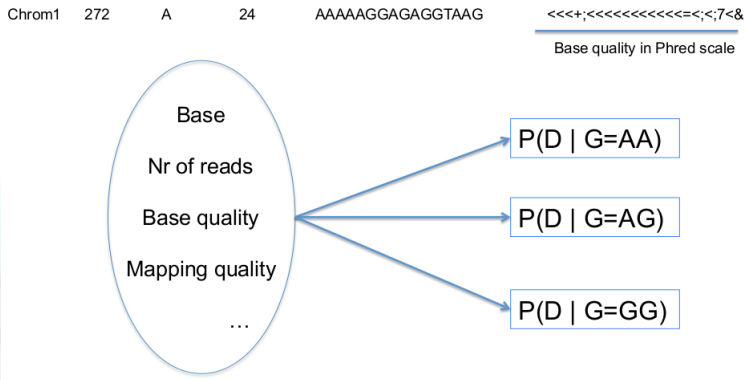
$$P(D|G = \{A_1, A_2, \dots, A_n\})$$

with

$A_i \in \{A, C, G, T\}$  and  $n$  being the ploidy

How many genotypes likelihoods do we need to calculate for each individual at each site?

# Genotype likelihoods



## Calculating genotype likelihoods

### Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

- $L_{A_j,i} = P(D|A_G = A_j)$
- $A_i \in \{A, C, G, T\}$
- $R$  is the depth (nr. of reads)
- $N$  is the ploidy (nr. of chromosomes)

Example:

A  
A  
A  
G

with all quality scores equal to 20 (in phred score)

$P(D|G = AC) = ?$

## Calculating genotype likelihoods

### Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

A  
A  
A  
G  
& Q=20

$$P(D|G = \{A, C\}) = \dots$$



## Calculating genotype likelihoods

### Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

A

A

A

G

& Q=20

$N = 2; i = 1; A_1 = A; A_2 = C$

$$P(D|G = \{A, C\}) = \left(\frac{L_{A,1}}{2} + \frac{L_{C,1}}{2}\right) \times \dots$$

What are  $L_{A,1}$  and  $L_{C,1}$ ?

## Calculating genotype likelihoods

## Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

A  
A  
A  
G  
& Q=20

$$L_{C,1} = \frac{\epsilon_1}{3}$$

$$L_{A,1} = 1 - \epsilon_1$$

$$P(D|G = \{A, C\}) = \left(\frac{1 - \epsilon_1}{2} + \frac{\epsilon_1}{6}\right) \times \dots$$

## Calculating genotype likelihoods

## Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

A

A

A

G

&amp; Q=20

$$L_{C,1} = \frac{\epsilon_1}{3}$$

$$L_{A,1} = 1 - \epsilon_1$$

$$P(D|G = \{A, C\}) = \left(\frac{1 - \epsilon_1}{2} + \frac{\epsilon_1}{6}\right) \times \left(\frac{1 - \epsilon_2}{2} + \frac{\epsilon_2}{6}\right) \times \left(\frac{1 - \epsilon_3}{2} + \frac{\epsilon_3}{6}\right) \times \frac{\epsilon_4}{3}$$

What are  $\epsilon_1, \epsilon_2, \dots$ ?

Genotype likelihoods

## Calculating genotype likelihoods

Genotype	Likelihood (log10)	
AA	-2.49	
<b>AC</b>	<b>-3.38</b>	
AG	-1.22	A
AT	-3.38	A
CC	-9.91	A
CG	-7.74	G
CT	-9.91	& $\epsilon = 0.01$
GG	-7.44	
GT	-7.74	
TT	-9.91	

## Genotype calling

Genotype	Likelihood (log10)
AA	-2.49
AC	-3.38
AG	-1.22
AT	-3.38
CC	-9.91
CG	-7.74
CT	-9.91
GG	-7.44
GT	-7.74
TT	-9.91

AAAG &  $\epsilon = 0.01$

What is the genotype here?

## Genotype calling

Genotype	Likelihood (log10)
AA	-2.49
AC	-3.38
<b>AG</b>	<b>-1.22</b>
AT	-3.38
CC	-9.91
CG	-7.74
CT	-9.91
GG	-7.44
GT	-7.74
TT	-9.91

AAAG &  $\epsilon = 0.01$

What is the genotype?

AG.

### Maximum Likelihood

The simplest genotype caller:  
choose the genotype with the  
highest likelihood.

## Major and minor alleles

### Likelihood function

$$\log P(D|G = A) = \sum_{i=1}^R \log L_{A_j,i}$$

AAAG &  $\epsilon = 0.01$

Allele	log-Likelihood
<b>A</b>	<b>-2.49</b>
C	-3.38
<b>G</b>	<b>-1.22</b>
T	-3.38

We can reduce the genotype space to 3 entries (from 10).

## Genotype likelihoods

AAAG & '5555' & A,G alleles

Genotype	log-Likelihood
AA	-5.73
AG	-2.80
GG	-17.12

Examples varying qualities and reads...



## Genotype likelihoods - example

AAAG & '5550' & A,G alleles

Genotype	log-Likelihood
----------	----------------

## Genotype likelihoods - example

AAAG & '5550' & A,G alleles

Genotype	log-Likelihood
AA	-4.58
AG	-2.81
GG	-17.14

## Genotype likelihoods - example

AAAG & '555K' & A,G alleles

Genotype	log-Likelihood
----------	----------------

## Genotype likelihoods - example

AAAG & '555K' & A,G alleles

Genotype	log-Likelihood
AA	-10.80
AG	-2.80
GG	-17.11

## Genotype likelihoods - example

AAAAAAAAAG & '555555550' & A,G alleles

Genotype	log-Likelihood
----------	----------------

## Genotype likelihoods - example

AAAAAAAAAG & '555555550' & A,G alleles

Genotype	log-Likelihood
AA	-4.64
AG	-7.01
GG	-51.37

# NGS data uncertainty

## Issue

There is a notable amount of statistical uncertainty in assigning individual genotypes depending on the number of reads and their quality. How can we deal with that when doing population genetics analysis (e.g. estimating genetic diversity)?

## Solutions

1. let's pretend we don't have such uncertainty

# NGS data uncertainty

## Issue

There is a notable amount of statistical uncertainty in assigning individual genotypes depending on the number of reads and their quality. How can we deal with that when doing population genetics analysis (e.g. estimating genetic diversity)?

## Solutions

1. let's pretend we don't have such uncertainty
2. genotype filtering
3. (a third way)



## Genotype likelihood ratio

$$\log_{10} \frac{L_{G(1)}}{L_{G(2)}} > t$$

i.e.  $t = 1$  meaning that the most likely genotype is 10 times more likely than the second most likely one

## Genotype posterior probability

AAAG &  $\epsilon = 0.01$  & A,G alleles

Genotype	Likelihood (log)	Prior	Posterior
AA	-5.73	1/3	0.05
AG	-2.80	1/3	0.95
GG	-17.12	1/3	0

Only call genotypes if the largest probability is above a certain threshold (e.g. 0.95).

Pros and cons?

## Genotype posterior probability

AAAG &  $\epsilon = 0.01$  & A,G alleles

Genotype	Likelihood (log)	Prior	Posterior
AA	-5.73	1/3	0.05
AG	-2.80	1/3	0.95
GG	-17.12	1/3	0

Only call genotypes if the largest probability is above a certain threshold (e.g. 0.95).

Pros and cons?

- Yes: genotype are called with higher **confidence**
- No: more **missing** data

## Exercise 1

Simulate NGS data and calculate genotype likelihoods and probabilities.

Assess the amount of uncertainty and data missingness.

nr of hetero individuals

assuming known genotypes with example of 4 samples

do not show real AA AG AG GG real

data A AAAG AGGG GG

unknown genotypes:

calculate by sampling distribution of nr of hetero from geno post

probs unif

then show expected value

genetic distances

pca

## Exercise 2

PCA with both methods

Alone: best PCA with filtering (perhaps later when introducing  
snp calling)



## Estimating allele frequencies

Assuming 2 alleles (A,G) with true allele frequency of 0.50

Sample	True genotype	Reads allele A	Read allele G
1	AA	7	0
2	AA	25	1
3	AG	5	3
4	AG	4	4
5	GG	0	2
6	GG	0	4

What is the simplest estimator of allele frequencies?

## Estimating allele frequencies

Assuming 2 alleles (A,G) with true allele frequency of 0.50

Sample	True genotype	Reads allele A	Read allele G
1	AA	7	0
2	AA	25	1
3	AG	5	3
4	AG	4	4
5	GG	0	2
6	GG	0	4
Total		41	14

$$\hat{f} = \frac{\sum_{i=1}^N n_{A,i}}{\sum_{i=1}^N (n_{A,i} + n_{G,i})}$$

$$\hat{f} = 0.75$$

What is wrong with this estimator?

## Estimating allele frequencies

Assuming 2 alleles (A,G) with true allele frequency of 0.50

Sample	True genotype	Reads allele A	Read allele G
1	AA	7	0
2	AA	25	1
3	AG	5	3
4	AG	4	4
5	GG	0	2
6	GG	0	4
Total		41	14

$$\hat{n}_A = \sum_{i=1}^N (1 - \epsilon)n_{A,i} + \epsilon n_{G,i} - \epsilon n_{A,i} - (1 - \epsilon)n_{G,i}$$

$$\hat{f} = 0.77$$

## Estimating allele frequencies

### Maximum Likelihood estimator

$$P(D|f) = \prod_{i=1}^N \sum_{g \in \{0,1,2\}} P(D|G = g)P(G = g|f)$$

## Estimating allele frequencies

### Maximum Likelihood estimator

$$P(D|f) = \prod_{i=1}^N \sum_{g \in \{0,1,2\}} P(D|G = g)P(G = g|f)$$

$P(D|G = g)$  is the genotype likelihood and  $P(G = g|f)$  is given by HWE (for instance).

In our previous example,  $\hat{f} = 0.46$  which is much closer to the true value than previous estimators.

# SNP calling

## Challenges

- If high levels of missing data, then genotypes can be lost.
- Rare variants are hard to detect.
- Trade off between false positive and false negative rates.

## How to call SNPs?

- If at least one heterozygous genotype has been called.
- If the estimated allele frequency is above a certain threshold.

## SNP calling

Call a SNP if

$$\hat{f} \geq t$$

where  $t$  can be the minimum sample allele frequency detectable (e.g.  $t = 1/2N$  with  $N$  diploids).

## Likelihood Ratio Test

A Likelihood Ratio Test (LRT) compares the goodness of fit between the null and the alternative model:

- Null model:  $f = 0$
- Alternative model:  $f \neq 0$

$$T = -2 \log \frac{L(f = 0)}{L(f = \hat{f}_{MLE})}$$

where  $T$  is  $\chi^2$  distributed with 1 degree of freedom.



## Exercise 3

Estimate allele frequencies and call SNPs  
calculate PCA with some filtering

expected value of proba being variable

expected value of nr snps



expected value of summary stats ( $\pi$ )

## Exercise - 4

calculate SFS

calculate summary stats 1 pop

alone: sliding windows

alone: 2 pops (2d-sfs and fst)

exp design with bigfoot

Thank you for your attention